

Interim Architecture Report

Surafel Yimam

Executive Summary

The Interim submission establishes the forensic infrastructure layer of the Automaton Auditor: a typed, parallelized, and sandboxed evidence collection system built on LangGraph.

This system is not a grading prompt wrapper. It is a structured digital investigative engine designed to:

- Objectively collect verifiable evidence from GitHub repositories
- Parse architectural structure using AST-based analysis
- Validate documentation claims against real code artifacts
- Preserve state integrity across parallel execution

The current implementation completes Phase 1 and Phase 2 of the required architecture:

- Typed State Management using Pydantic + reducers
- Parallel Detective Layer with fan-out/fan-in orchestration
- Secure Git Sandboxing and AST-based structural parsing
- Structured Evidence objects suitable for judicial synthesis

Judicial personas and deterministic synthesis will be implemented in the Final submission.

1. Why Pydantic Over Plain Dictionaries

In a parallel multi-agent system, untyped dictionaries are architectural liabilities.

Problems With Dict-Based State

- Silent overwrites in parallel execution
- Schema drift across nodes
- No validation guarantees
- Impossible to enforce reducer semantics

Architectural Decision

I implemented:

- `Evidence` as a strict Pydantic BaseModel
- `JudicialOpinion` and `AuditReport` as forward-compatible BaseModels
- `AgentState` as a TypedDict with Annotated reducers

Reducers used:

- `operator.add` for list accumulation
- `operator.ior` for dictionary merges

This guarantees:

- Parallel branches cannot overwrite each other
- Evidence accumulates deterministically
- Future Judicial Layer integration will not require refactoring

This design prevents “dict soup” and enforces structural rigor from the foundation upward.

2. AST-Based Graph Verification (Why Regex Is Insufficient)

A naive implementation would scan for the string “StateGraph”.

That approach is brittle and fails under:

- Variable aliasing
- Refactoring
- Multi-line formatting
- Indirect instantiation

Instead, I use Python’s built-in `ast` module to:

- Parse `src/graph.py` into an Abstract Syntax Tree
- Identify actual class instantiations of `StateGraph`
- Detect method calls to `add_edge` and `add_conditional_edges`
- Determine whether parallel fan-out patterns exist

This allows structural verification rather than surface-level string matching.

This is critical for scoring under “Deep AST Parsing” in Forensic Accuracy.

3. Secure Tool Engineering & Sandboxing Strategy

The system clones unknown repositories. This is inherently dangerous.

Security Threat Model

If implemented carelessly:

- Code could execute during clone hooks
- Shell injection could occur
- Files could overwrite working directory
- Malicious repo names could manipulate shell commands

Defensive Architecture

The RepoInvestigator:

- Uses `tempfile.TemporaryDirectory()`
- Uses `subprocess.run()` (never `os.system`)
- Captures `stdout/stderr`
- Validates return codes
- Never clones into live working directory

This ensures:

- Isolation
- Controlled execution
- Clean teardown

- No persistence of foreign code

This design satisfies Safe Tool Engineering requirements.

4. Parallel Orchestration (Fan-Out / Fan-In)

The LangGraph architecture is explicitly parallel.

Flow:

START

→ ContextBuilder
→ (RepoInvestigator || DocAnalyst)
→ EvidenceAggregator
→ END

Key properties:

- Detectives execute concurrently
- Evidence merges via reducers
- No node overwrites state
- Architecture is extensible to additional branches

This establishes the foundational swarm topology required for Week 2.

5. Theoretical Depth Verification (DocAnalyst Design)

The DocAnalyst does not blindly grep keywords.

It:

1. Extracts PDF text
2. Chunks it
3. Searches for critical architectural terms:
 - Dialectical Synthesis
 - Fan-In / Fan-Out
 - Metacognition
 - State Synchronization
4. Evaluates contextual depth (buzzword vs explanation)

This prepares the system for cross-referencing report claims against repository artifacts in the Final phase.

6. Known Gaps & Forward Plan

This Interim submission intentionally focuses on forensic infrastructure.

Not yet implemented:

- Judicial personas (Prosecutor / Defense / Tech Lead)
- Deterministic Chief Justice synthesis
- Conflict resolution rules
- Markdown AuditReport rendering
- Conditional error edges
- VisionInspector execution (stub only)

Plan for Final Submission

- Implement parallel Judge fan-out
- Bind Judges to structured Pydantic outputs
- Implement deterministic synthesis rules:
 - Security Override
 - Fact Supremacy
 - Functionality Weight
 - Variance Re-evaluation
- Serialize AuditReport to Markdown
- Integrate LangSmith tracing

7. Architectural Philosophy

This system is intentionally designed as infrastructure, not a grading prompt.

It reflects three core principles:

1. Evidence before opinion
2. Structure before scale
3. Determinism before delegation

The Interim submission establishes the investigative department of a Digital Courtroom.

The Final submission will implement the dialectical bench and synthesis engine.

Conclusion

The Automaton Auditor Interim version achieves:

- Typed state rigor
- Secure repository analysis
- AST-level structural verification
- Parallel orchestration
- Structured evidence accumulation

It is architecturally sound, extensible, and production-oriented.

This is not a toy agent.

It is governance infrastructure.